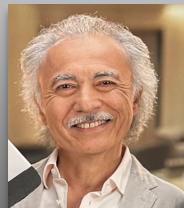


Semantic Search from the Command Line

***Learning about Embeddings/Similarity
for Java Developers***

Gran Sasso

NYJavaSIG



Frank Greco

@frankgreco

[linkedin.com/in/frankdgreco](https://www.linkedin.com/in/frankdgreco)

Chair
NYJavaSIG - NY Java User Group
@nyjavasig



Topics

- *(5) Problem Description*
- *(20) What are Embeddings and What is Similarity*
- *(5) LangChain4j Overview*
- *(5) Get Embeddings - demo*
- *(5) Compare Embeddings - demo*
- *(10) Vector Stores and Chunking*
- *(10) LangChain4j Ingestion - demo*
- *(5) Vector Search and Filters*
- *(5) Semantic Search - demo*
- *(5) Q/A*

Goals

Search for similar text content on your filesystem.

- *Need info, but can't remember filename or exact text*
- *Can't come up with a regexp-able string as search target*
- *Need to understand embeddings / "similarity" for GenAI project*
 - *Venkat did emphasize learning at D2N*
- *You are a CLI person*
- *Desperately need a distraction from following the news...*

FIRST...

**WHAT IS THE
PROBLEM?**

Problem to solve

You need to recall some important information.

*You recall that information in a document on your filesystem,
but you don't remember the filename, the exact text.
And you can't come up with a useful regexp as a target.*

```
grep - Standard GNU grep utility
egrep - Extended grep (equivalent to grep -E)
fgrep - Fast grep for fixed strings (equivalent to grep -F)
rgrep - Recursive grep (equivalent to grep -r)
pgrep - Process grep for searching processes
zgrep - Grep for gzip-compressed files
bzgrep - Grep for bzip2-compressed files
xzgrep - Grep for xz-compressed files
ngrep - Network packet grep
agrep - Approximate grep for fuzzy matching
sgrep - Structured grep for XML/JSON
ugrep - Ultra grep with TUI and fuzzy search
ripgrep (rg) - Fast grep respecting gitignore
ggrep - GNU grep on non-GNU systems
vgrep - Visual grep with interactive TUI
hgrep - Grep with syntax highlighting
cgrep - Context-aware grep for source code
gitgrep - Grep in git repositories
lcfgrep - Grep with line context
tgrep - Grep with directory tree view
mgrep - Multi-pattern grep
wgrep - Word grep
jgrep - JSON-aware grep
xmlgrep - XML-aware grep
pygrep - Python-specific grep
jsgrep - JavaScript-specific grep
phpgrep - PHP-specific grep
gogrep - Go-specific grep
rggrep - Regex grep with advanced features
fngrep - Filename grep
lgrep - Line-oriented grep
dgrep - Disk grep for searching file systems
igrep - Interactive grep
kgrep - Kernel grep for searching Linux kernel source
webgrep - Web page grep
csvgrep - CSV-aware grep
sqlgrep - SQL-aware grep
bingrep - Binary file grep
```

The Grep Family

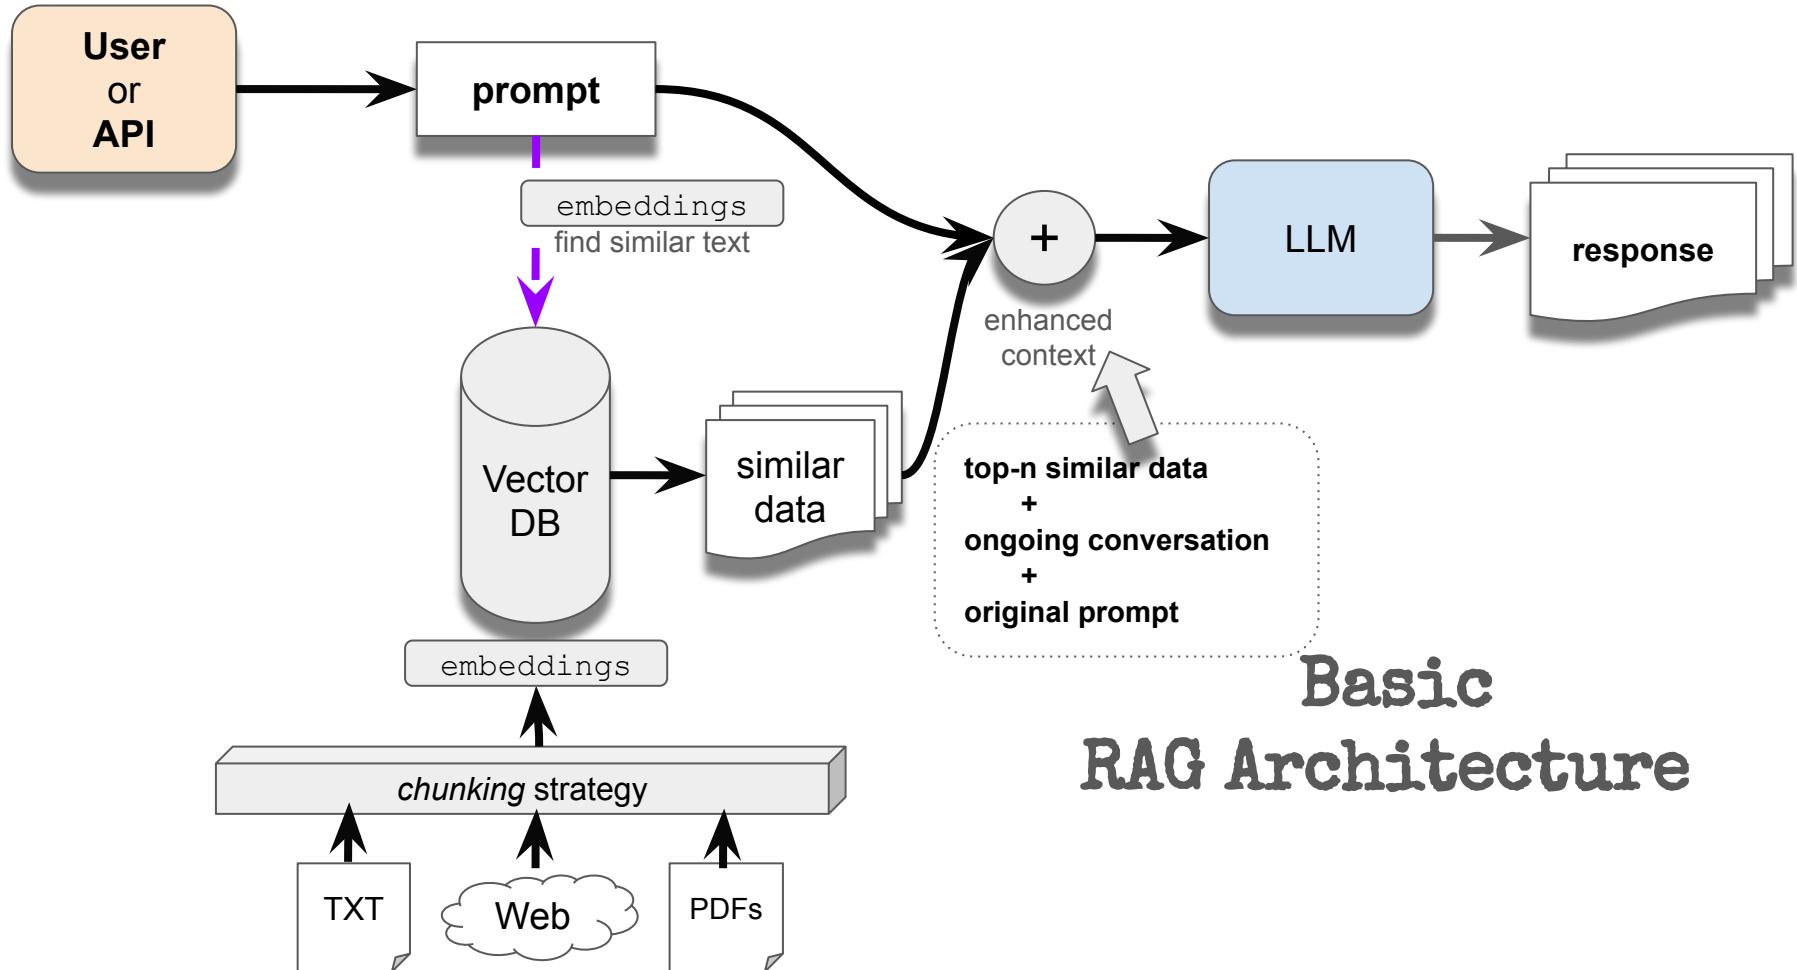
Some have regexp and “fuzzy” matching.

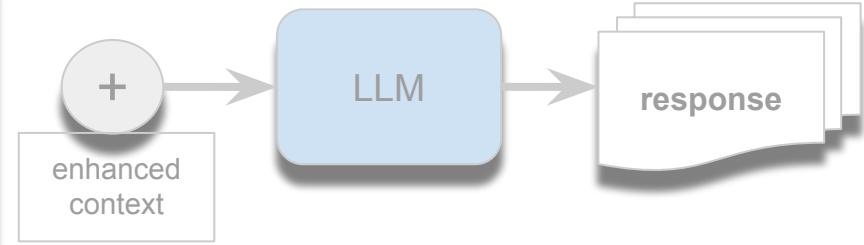
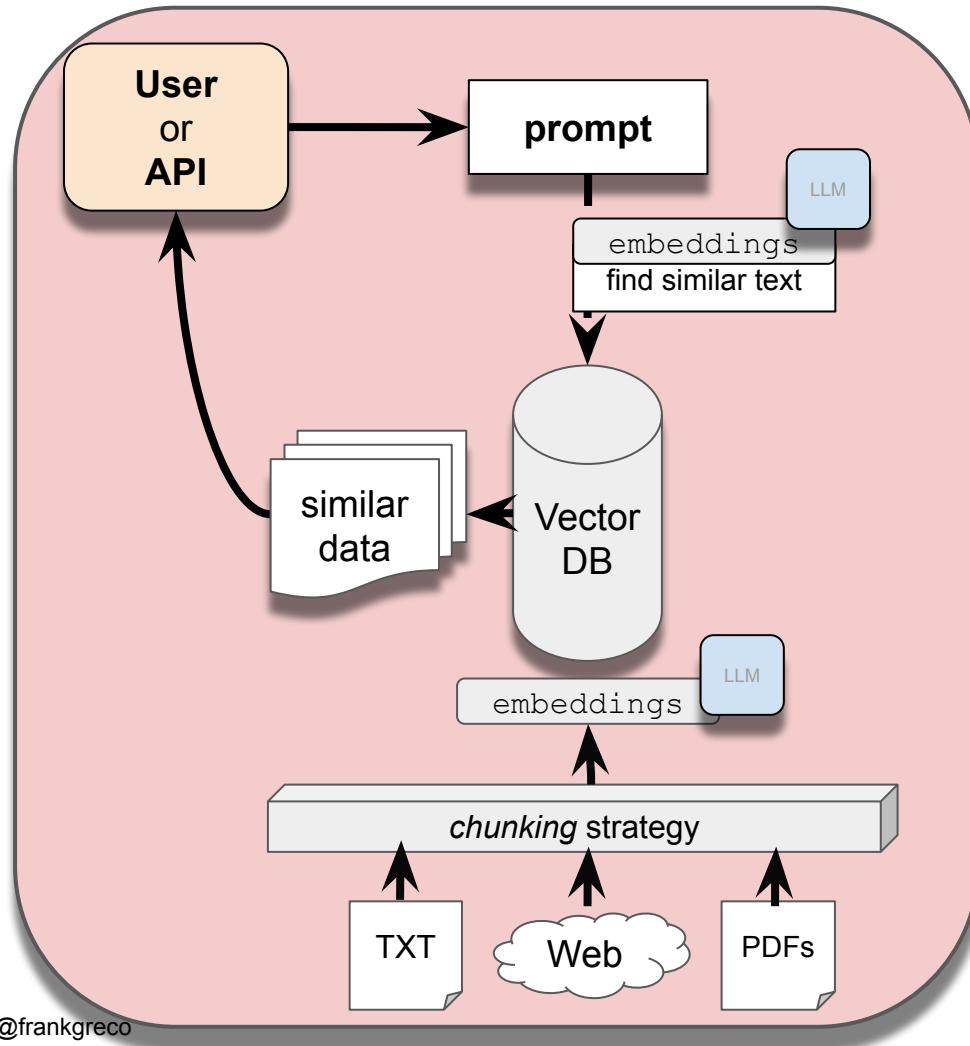
Regexp doesn’t account for ***similar*** matches.
Fuzzy matches have high false positives.

Arguable whether *all* of these are actually necessary.
But I digress...



WASN'T THERE SOMETHING ABOUT A
BASIC GENAI RAG SYSTEM
THAT DETERMINED SIMILARITY?



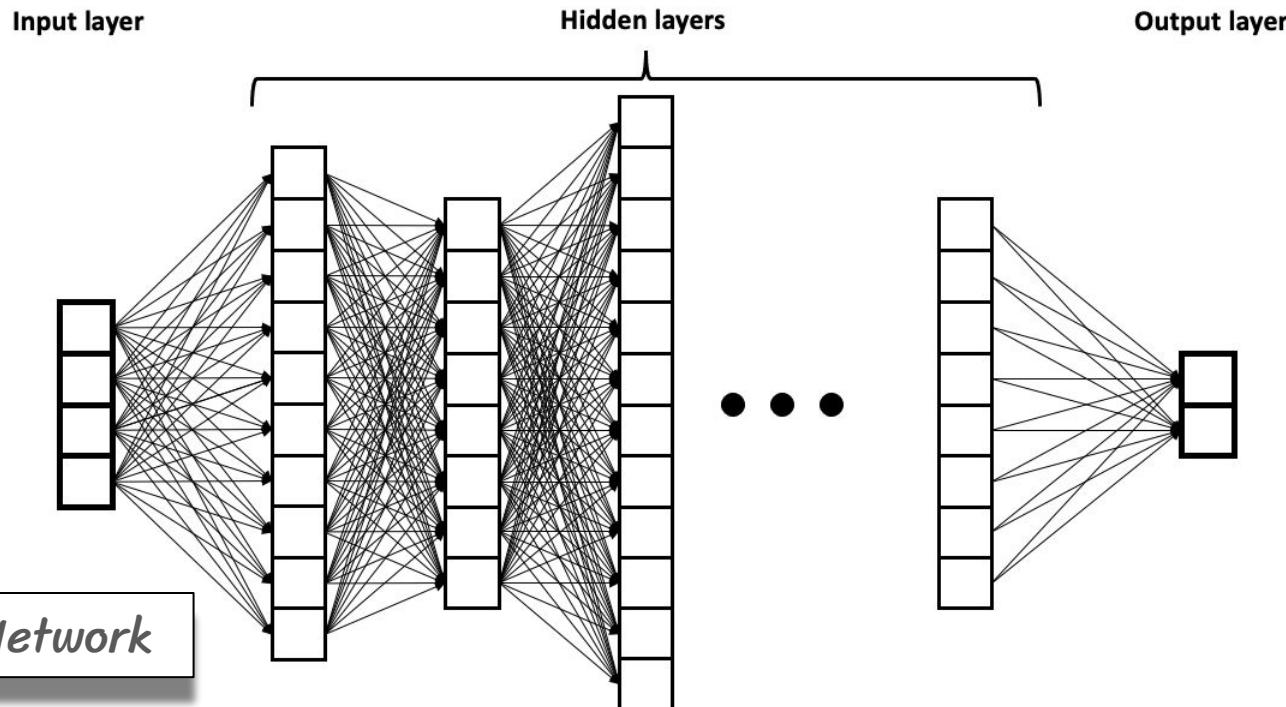


Basic RAG Architecture

EMBEDDING VECTORS AND SIMILARITY

©Frank Greco

*In AI applications,
Embeddings are used to determine **Similarity** .*



Similarity Applications

Text

Semantic search, sentiment analysis, language understanding, anomaly detection, grouping support issues

Images

Similar images, style detection, clustering, visual search, classify video scenes

Audio

Music recommendation, sound pattern detection



Let's talk about Text Embeddings

Text Embeddings

Fixed-length, numerical vectors representing words, sentences, paragraphs, or even entire documents.

“Embedding vector”

Using Java/LangChain4j

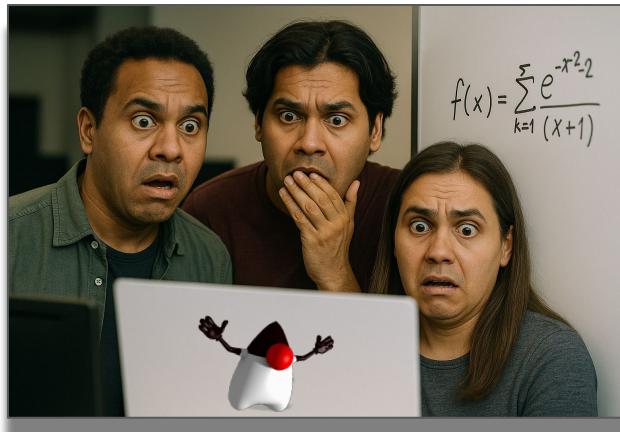
```
Embedding vector = [-0.005175813, -0.013680333, 0.018862674,  
0.024423573, 0.024097228, -0.0039813947, -0.037333734, -0.003062737,  
-0.033208746, 3.4674033E-4, 0.0653732, -0.04566203, 0.002189767,  
0.013823925, 0.0059884093, 0.010958626, -0.0035310402, -0.037960313,  
-0.021460371, 0.051118497, 0.01725706, -0.064015605, 0.015742827,  
0.022517724, 0.016826289, 0.03213834, ...]
```

*Ewww... I'm afraid of math!!
There are no semi-colons or curly braces!*



Text Embeddings are Useful!

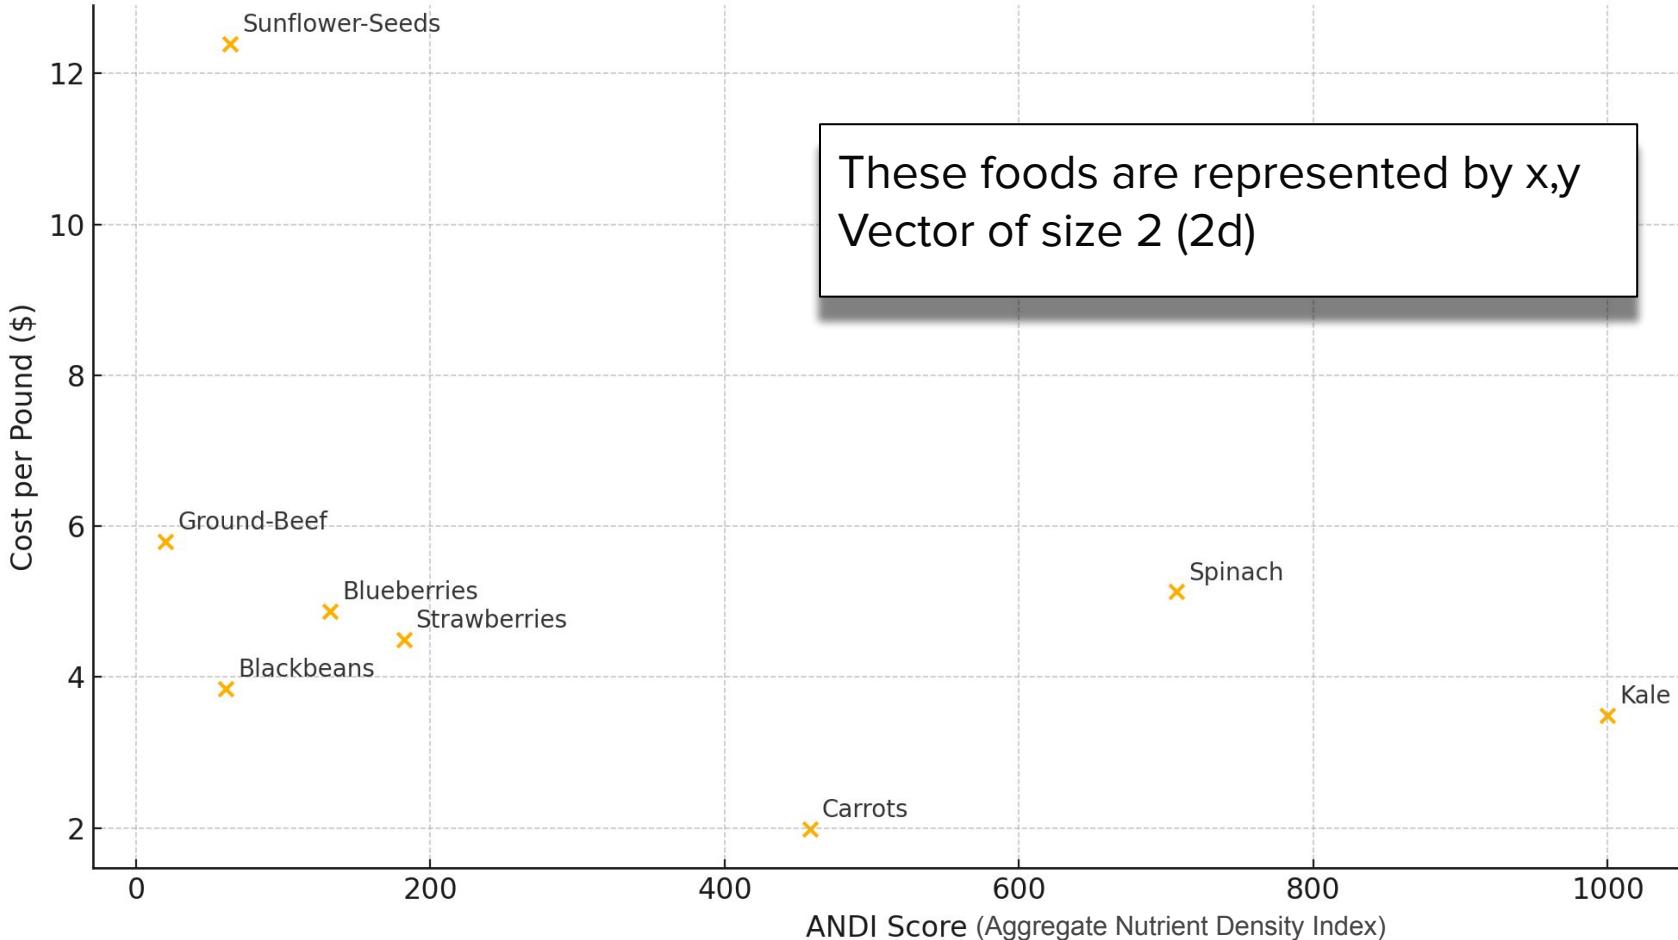
Text embeddings are used to extract semantic meaning and calculate similarity.



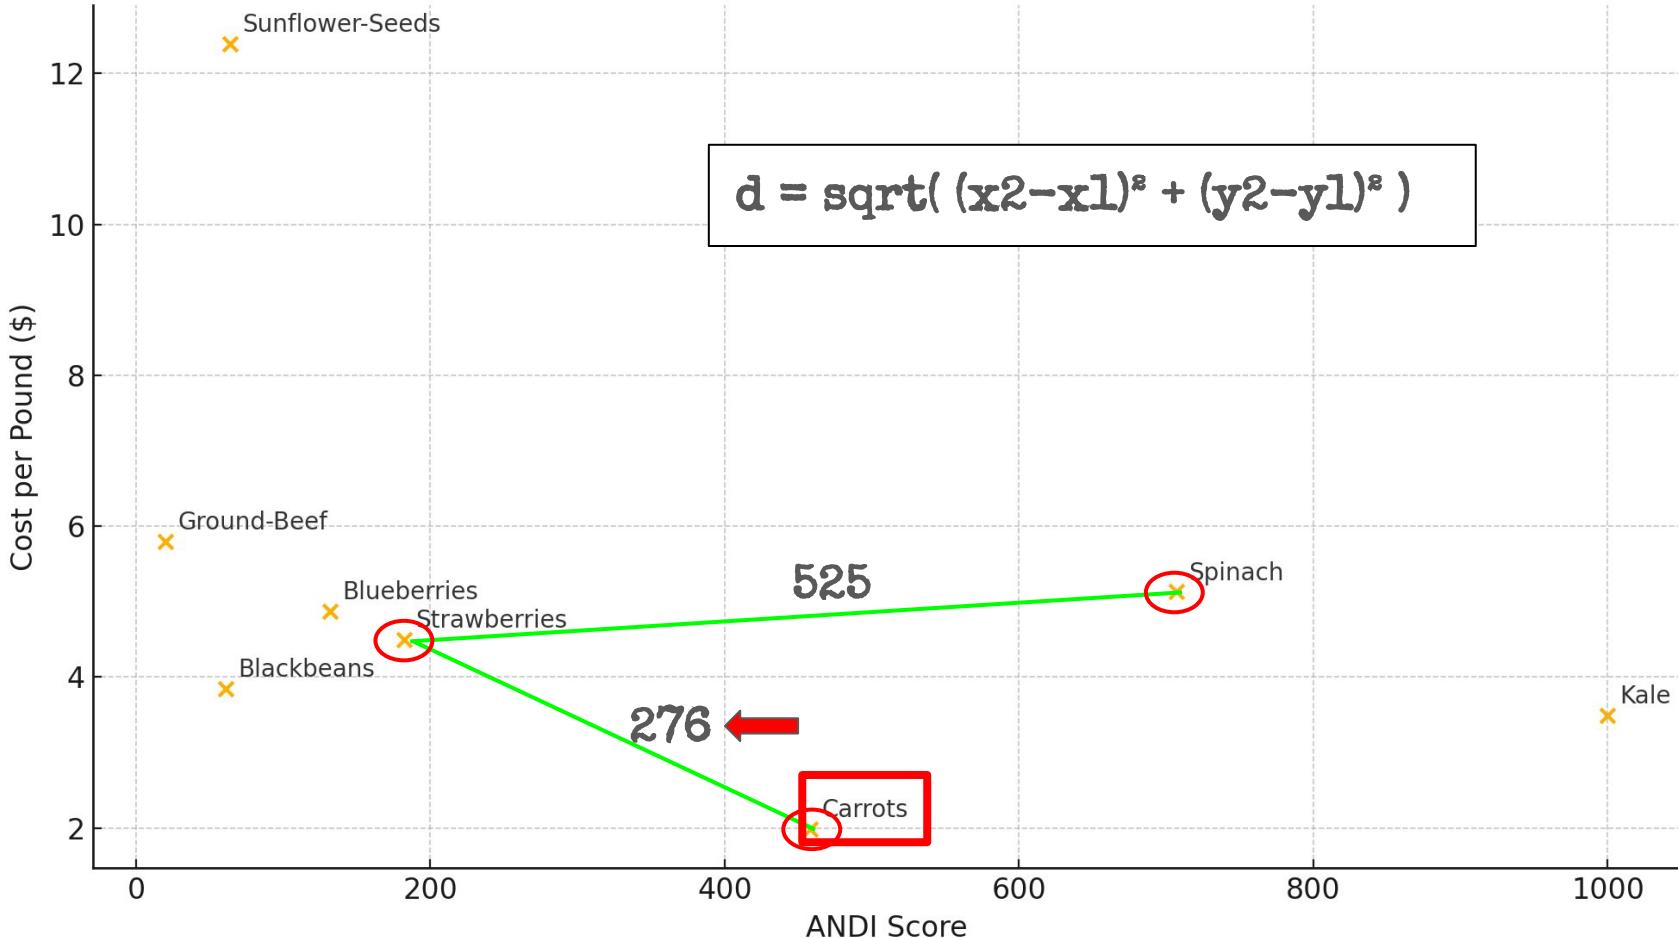
Wait a minute,

What exactly is “Similarity”?

Nutrient Density (ANDI) vs Cost per Pound



Nutrient Density (ANDI) vs Cost per Pound





1



2



3

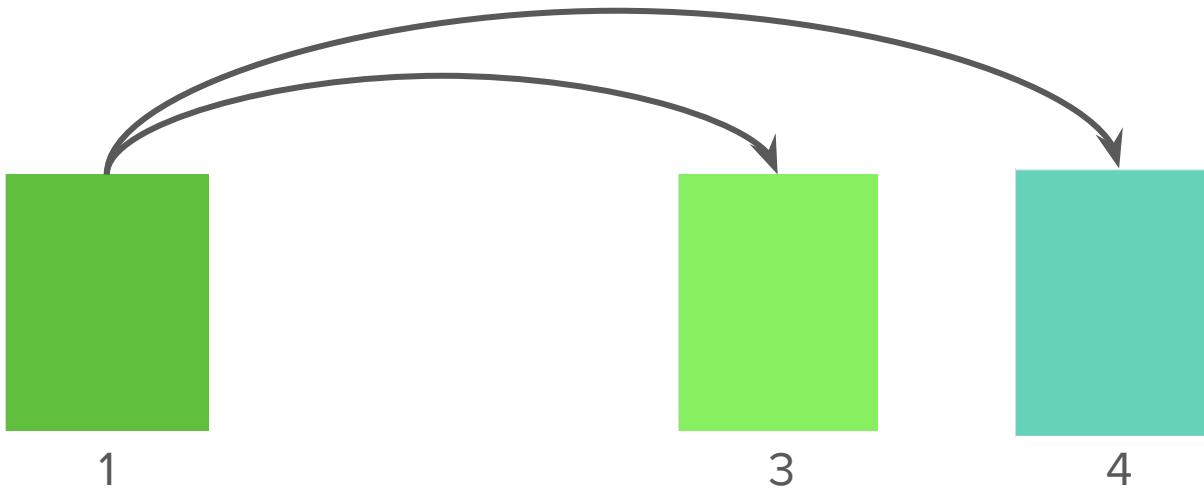


4



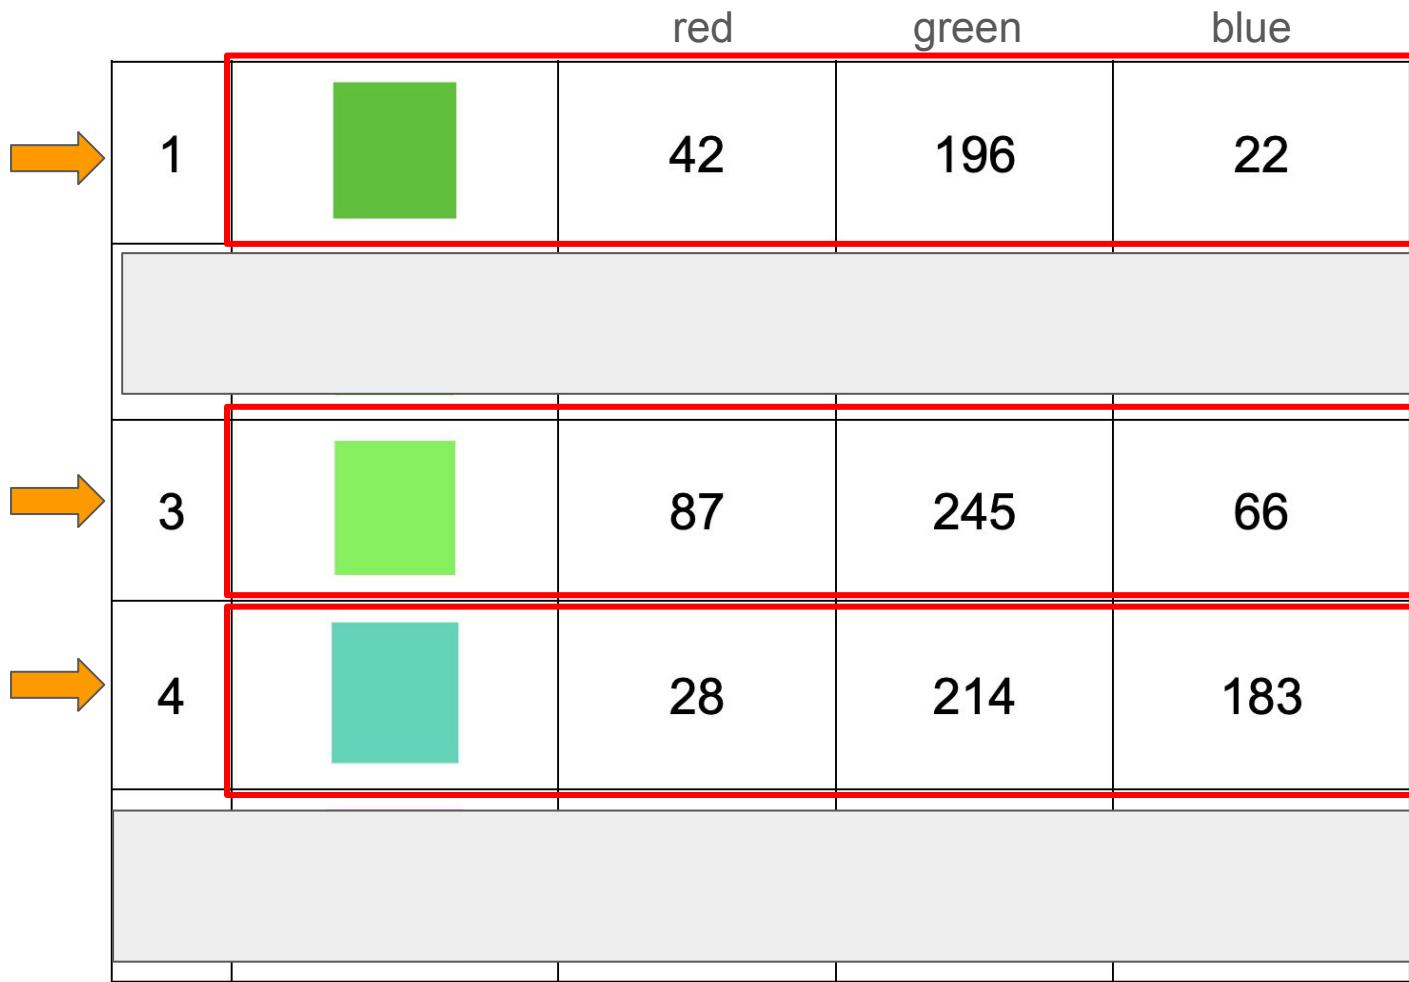
5

Color similarity



Is Color-3 or Color-4 more similar to Color-1?

		red	green	blue
1		42	196	22
3		87	245	66
4		28	214	183

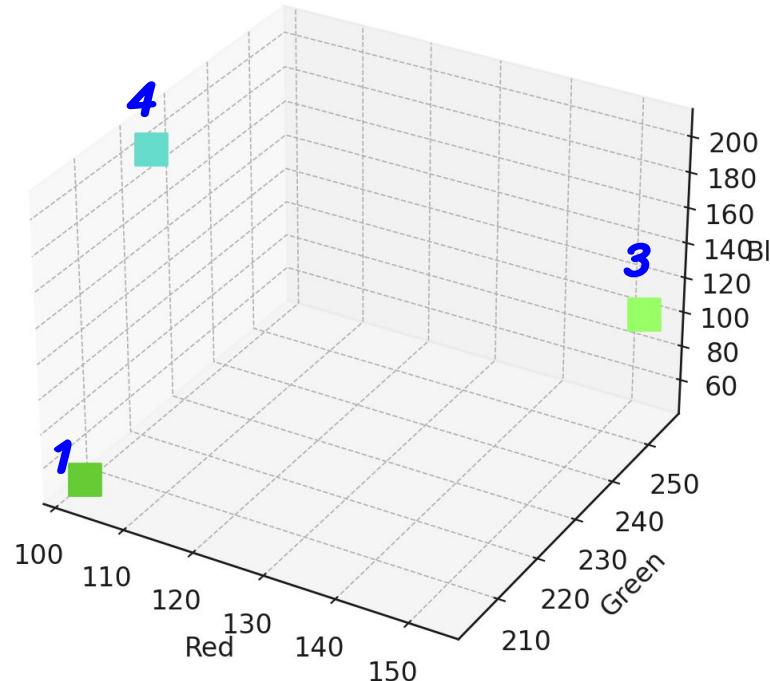


The diagram illustrates three rows of color data, likely representing different colors in a digital image or print process. Each row consists of five columns: a row index (1, 3, or 4), a column index (empty), and three numerical values (red, green, blue). The first row contains a solid green square, the second a solid light green square, and the third a solid teal square. Each square is enclosed in a red rectangular border. To the left of each row, an orange arrow points towards it. The numerical values are as follows:

Row	Column	Red	Green	Blue
1		42	196	22
3		87	245	66
4		28	214	183

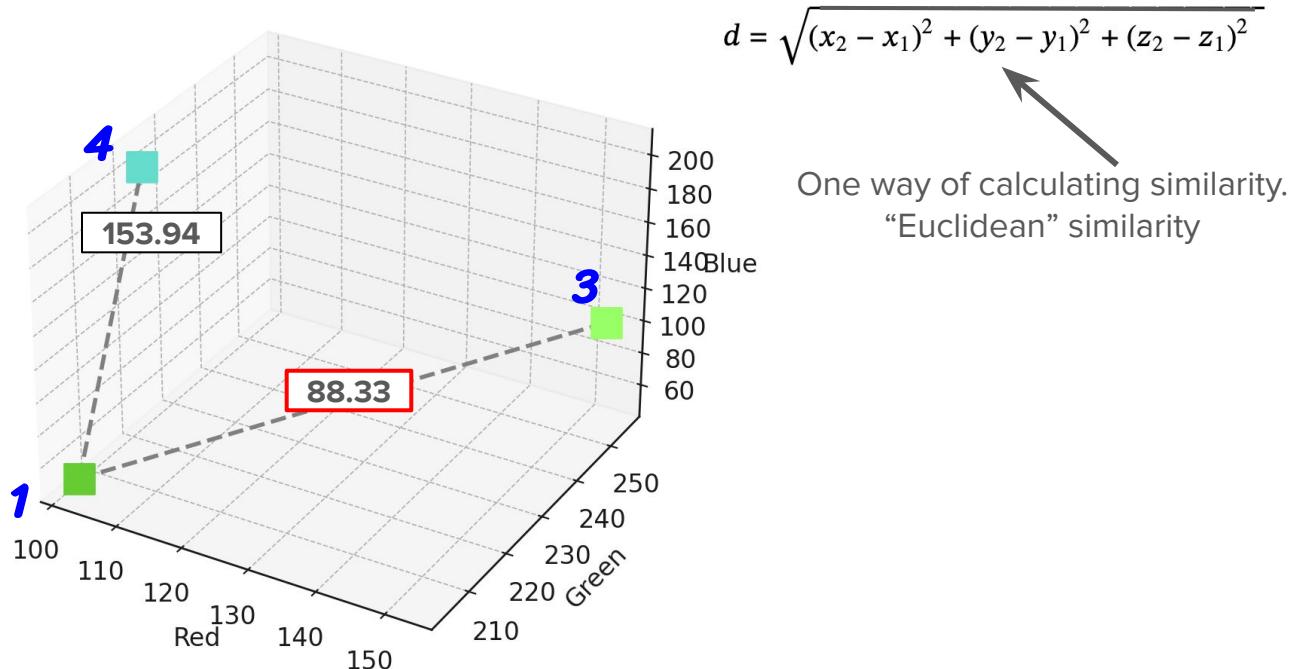
These colors are represented by R,G,B
Vector of size 3 (3d)

3D RGB Color Space (Points 1, 3, 4 Only)



These colors are represented by R,G,B
Vector of size 3 (3d)

3D RGB Color Space with Connecting Lines (Points 1, 3, 4)



INTERESTING!

BUT WHAT ABOUT TEXT SIMILARITY?

*Instead of 2 or 3 dimensions,
Text (String) characteristics have a lot!*

Num of dimensions related to size of LLM's internal neural network.

Typical # dimensions for vector: 384, 768, 1536, 2048, etc.

*Since neural networks are not all the same size,
embedding vectors are not all the same size.*

“Embedding” originates from mathematics.

Mapping from one space to another.

Food -> (2d) nutrition and cost

Color -> (3d) R, G and B

Text Strings -> (1,536d) many characteristics...

There are dozens of embedding services and libraries to use.

Btw, don't mix embedding vectors from different sources!

Embeddings Service

```
import ...          // LangChain4j imports

public class GetEmbedding {

    public static void main(String[] argv) {
        String apiKey = System.getenv("OPENAI_API_KEY");

        EmbeddingModel model = OpenAiEmbeddingModel.builder()
            .apiKey(apiKey)
            .modelName(TEXT_EMBEDDING_3_SMALL)
            .build();

        Response<Embedding> response = model.embed("One fish, two fish, red fish, blue fish");
        Embedding embedding = response.content();

        System.out.println(embedding);
    }
}
```

OpenAI's Embedding Service

1536 elements: Embedding { vector = [-0.005175813, -0.013680333, 0.018862674, 0.024423573,
0.024097228, -0.0039813947, -0.037333734, -0.003062737, -0.033208746, 3.4674033E-4,
0.0653732, -0.04566203, 0.002189767, 0.013823925, 0.0059884093, 0.010958626,
-0.0035310402, -0.037960313, -0.021460371, 0.051118497, 0.01725706, -0.064015605,
0.015742827, 0.022517724, 0.016826289, 0.03213834, 6.065916E-4, 0.03098961, 0.021512585,
0.018875727, 0.073779814, -0.016943771, 0.0035473574, -0.016434675, -0.057227653,
0.004944109, 0.0071665104, -0.024397464, 0.015181515, -0.0010826456, 0.018901834,
0.048690498, 0.054825764, -0.013641172, 0.06960261, 0.020794628, -0.043286245,
-0.028666042, -0.005857872, 0.06552984, -0.040127236, 0.08552819, -0.054877978,
0.037620917, -0.0023317267, -0.008654638, -0.049499832, 0.037829775, 0.024045013,
-0.058271956, 0.054773547, 0.032477736, 0.0038214861, 0.010025282, -0.020246372,
0.040988784, -0.047698416, -0.017792266, 0.052084476, 0.018458007, -0.016056117,
0.020481339, 0.014554935, 0.013099441, 0.044878803, 0.025063206, -0.043703966,
-0.034096405, -6.1026297E-4, 0.054773547, -0.01366728, 0.017374545, 0.018027233,
-0.05540013, -0.023849208, -0.03161619, 0.0140197305, -0.0344358, -0.025520088,
-0.013758656, -0.0048527326, -0.04532263, -0.088400014, -0.03955287, -0.001202577,
-0.022857122, 0.024084175, 0.015077085, -0.034383588, 0.025781162, 0.019658953,
-0.04106711, -0.041328184, -0.020559661, 0.0057208072, 0.011500357, 0.03310432,
-0.008367456, -0.0011772853, -0.005985146, -0.01847106, 0.005883979, 0.015155408,
-9.129469E-4, -0.00326507, 0.054721333, 0.008387037, -0.029266514, -0.0054303613,
-0.023314003, 0.022491617, -0.027047377, 0.046314716, -0.031433437, -0.029240407,
0.0093987025, -0.010423422, -0.032608274, 0.0029240407, 0.00829566, -0.001980907,
0.021930305, -0.017178738, -0.008256499, -0.005567426, 0.024567163, -0.042842418, ...

```
1536 elements: Embedding { vector = [-0.005175813, -0.013680333, 0.018862674, 0.024423573,  
0.024097228, -0.0039813947, -0.037333734, -0.003062737, -0.033208746, 3.4674033E-4,  
0.0653732, -0.04566203, 0.002189767, 0.013823925, 0.0059884093, 0.010958626,  
-0.0035310402, -0.037960313, -0.021460371, 0.051118497, 0.01725706, -0.064015605,  
0.015742827, 0.022517724, 0.016826289, 0.03213834, 6.065916E-4, 0.03098961, 0.021512585,  
0.018875727, 0.073779814, -0.016943771, 0.0035473574 -0.016434675 -0.057227653  
0.004944109, 0.0071665104, -0.024397464, 0.  
0.048690498, 0.054825764, -0.013641172, 0.  
-0.028666042, -0.005857872, 0.06552984, -0.  
0.037620917, -0.0023317267, -0.008654638,  
-0.058271956, 0.054773547, 0.032477736, 0.  
0.040988784, -0.047698416, -0.017792266, 0.  
0.020481339, 0.014554935, 0.013099441, 0.0  
-0.034096405, -6.1026297E-4, 0.054773547, -  
-0.05540013, -0.023849208, -0.03161619, 0.0140197305, -0.0344358, -0.025520088,  
-0.013758656, -0.0048527326, -0.04532263, -0.088400014, -0.03955287, -0.001202577,  
-0.022857122, 0.024084175, 0.015077085, -0.034383588, 0.025781162, 0.019658953,  
-0.04106711, -0.041328184, -0.020559661, 0.0057208072, 0.011500357, 0.03310432,  
-0.008367456, -0.0011772853, -0.005985146, -0.01847106, 0.005883979, 0.015155408,  
-9.129469E-4, -0.00326507, 0.054721333, 0.008387037, -0.029266514, -0.0054303613,  
-0.023314003, 0.022491617, -0.027047377, 0.046314716, -0.031433437, -0.029240407,  
0.0093987025, -0.010423422, -0.032608274, 0.0029240407, 0.00829566, -0.001980907,  
0.021930305, -0.017178738, -0.008256499, -0.005567426, 0.024567163, -0.042842418, ...
```

You rarely look at these values.

Just understand that these values represent how one embedding algorithm views the characteristics of a string.

Each individual element isn't really meaningful by itself, but the vector represents the semantics of the string.

(Very) loosely speaking, each element is a string attribute.

HOW DO I CALCULATE
TEXT SIMILARITY??

There are several ways of calculating similarity of 2 points in N-dimensional space.

Probably more than a dozen...

Each algorithm has its advantages and disadvantages for different types of data.

Euclidean Similarity

$$\text{distance } (\mathbf{b}, \mathbf{a}) = \sqrt{\sum_{i=1}^n (b_i - a_i)^2}$$

Euclidean distance between two vectors

Similarity of two vectors
Closer to 0 implies similarity

17 types of similarity and dissimilarity measures used in data science.

The following article explains various methods for computing distances and showing their instances in our daily lives. Additionally, it will introduce you to the pydist2 package.



Mahmoud Harmouch · Follow

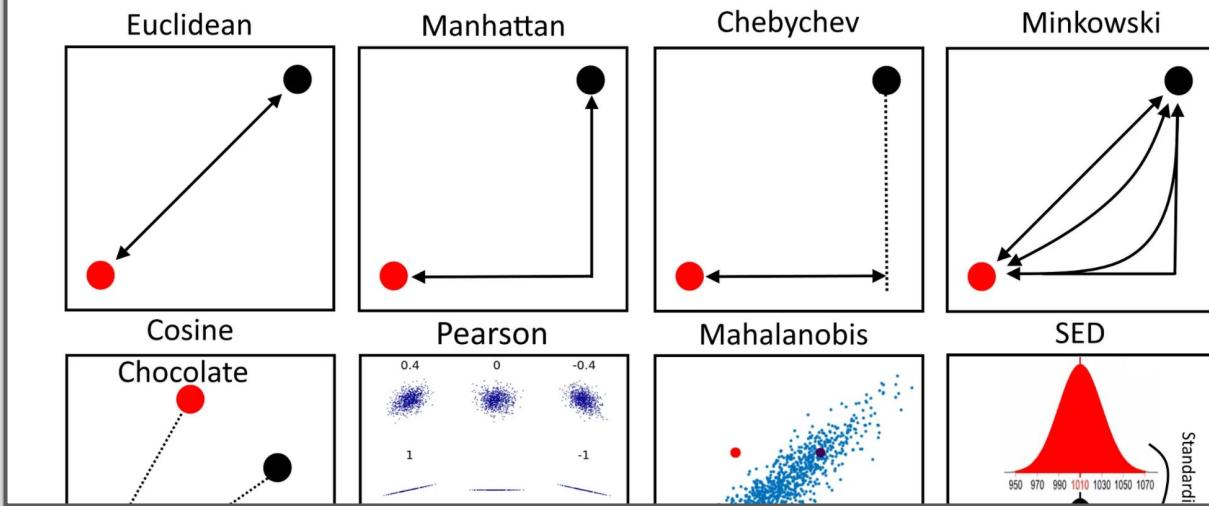
Published in Towards Data Science · 25 min read · Mar 13, 2021



2.7K



17



Cosine Similarity

Machine friendly since it uses matrices!

$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

Cosine of the angle between two vectors

Similarity of two vectors [-1, +1]

LANGCHAIN4J OVERVIEW

©Frank Greco

langchain4j / langchain4j

Type ⌘ to search

Code Issues 220 Pull requests 64 Discussions Actions Projects Security Insights

langchain4j Public

Watch 72 Fork 775 Star 4k

main 23 Branches 35 Tags Go to file Add file Code

langchain4j re #725: PR 5: Updated documentation. (#1473) ✓ fdba052 · 7 hours ago 1,014 Commits

.devcontainer Add Dev Container support (#337) 7 months ago

.github Update pull_request_template.md 2 weeks ago

.mvn(wrapper) Correctly configure Maven wrapper (#348) 7 months ago

code-execution-engines changed version to 0.33.0-SNAPSHOT last week

docker/ollama feat : create llama3 model image (#1083) 2 months ago

docs re #725: PR 5: Updated document 7 months ago

document-loaders changed version to 0.33.0-SNAPSHOT last week

document-parsers changed version to 0.33.0-SNAPSHOT last week

embedding-store-filter-parsers/langchain4j-... changed version to 0.33.0-SNAPSHOT last week

experimental/langchain4j-experimental-sql changed version to 0.33.0-SNAPSHOT last week

langchain4j-anthropic changed version to 0.33.0-SNAPSHOT last week

langchain4j-azure-ai-search Feat(#1383): mutualise EmbeddingMatches handling (#1... last week

langchain4j-azure-cosmos-mongo-vcore changed version to 0.33.0-SNAPSHOT last week

langchain4j-azure-cosmos-nosql changed version to 0.33.0-SNAPSHOT last week

langchain4j-azure-open-ai changed version to 0.33.0-SNAPSHOT last week

About

Java version of LangChain

docs.langchain4j.dev

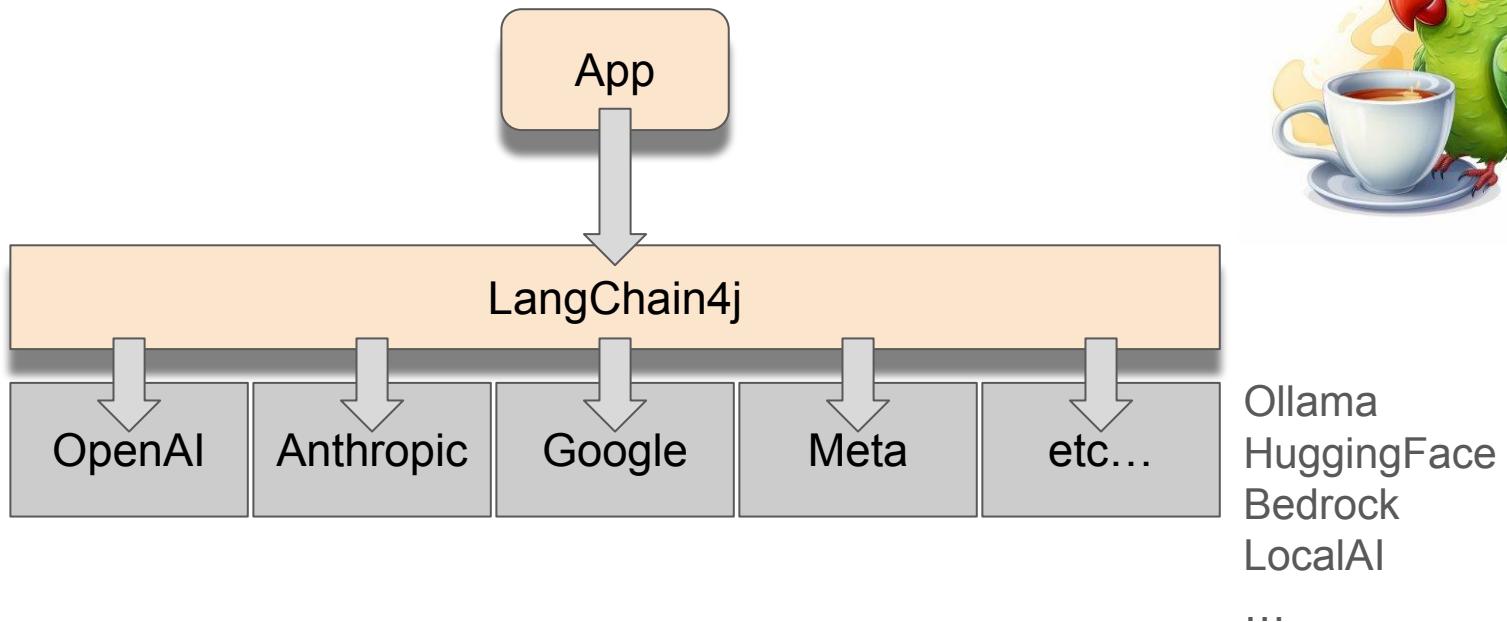
java embeddings gemini openai chroma llama gpt pinecone onnx weavite huggingface milvus vector-database openai-api chatgpt langchain anthropic pgvector ollama

<https://github.com/langchain4j/langchain4j>

Activity 4k stars 72 watching 775 forks Report repository

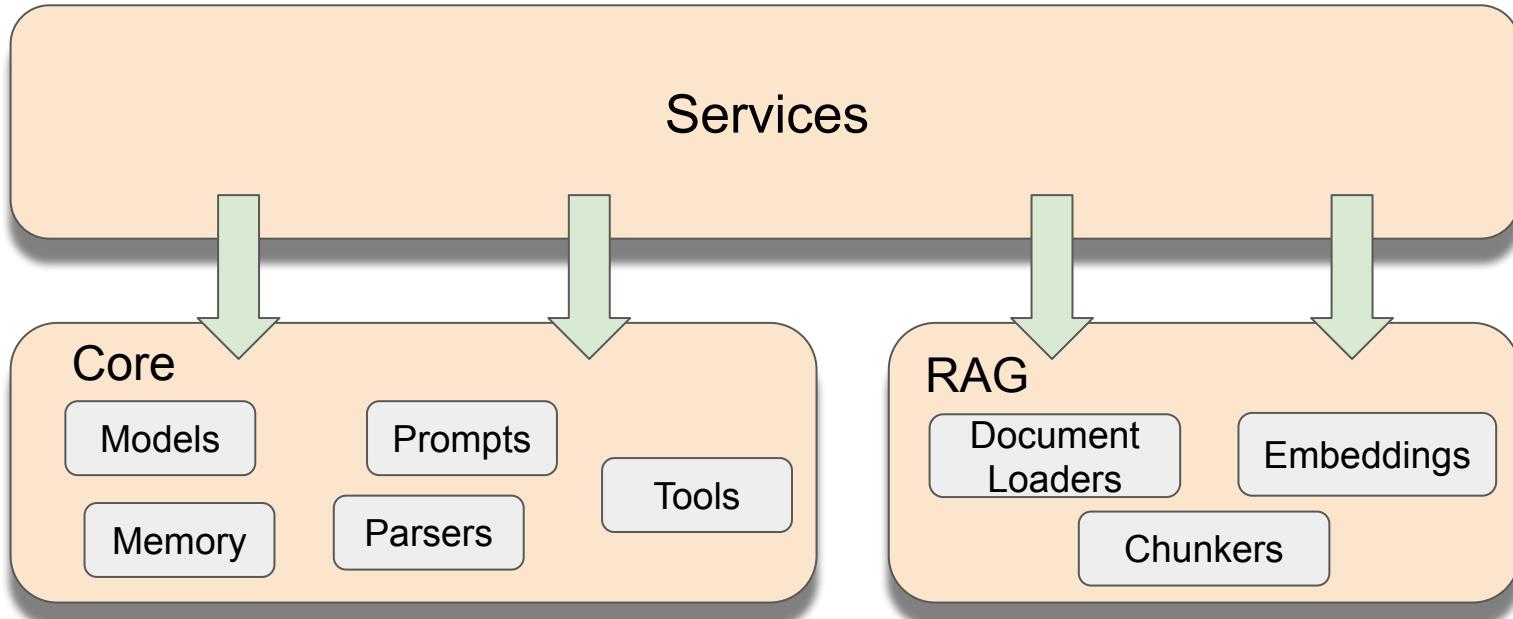
Releases 34

0.32.0 Latest 2 weeks ago + 33 releases



Provides an abstraction for LLMs and associated tools.

LangChain4J Components



Simple Completion

```
public class SimpleChatCompletion {  
    public static void main(String[] args) throws IOException {  
  
        ChatModel cmodel = OpenAiChatModel.builder()  
            .apiKey(System.getenv("OPENAI_API_KEY"))  
            .modelName(OpenAiChatModelName.GPT_4_O)  
            .timeout(Duration.ofSeconds(30))  
            .maxTokens(1024)  
            .build();  
  
        List<ChatMessage> messages = new ArrayList<>();  
  
        UserMessage usermsg = new UserMessage("How should I learn Java?");  
        messages.add(usermsg);  
  
        Response<AiMessage> answer = cmodel.chat(messages);  
  
        System.out.println(answer.content().text());  
    }  
}
```



Simple Completion

```
public class HelloWorldGoogle {  
    public static void main(String[] argv) {  
        String apiKey = System.getenv("GOOGLE_API_KEY");  
  
        ChatModel model = GoogleAiGeminiChatModel.builder()  
            .apiKey(apiKey)  
            .modelName("gemini-1.5-flash") ← Gemini  
            .build();  
  
        List<ChatMessage> messages = new ArrayList<>();  
        UserMessage usrmsg = UserMessage.from("Why should I learn Java.");  
        messages.add(usrmsg);  
  
        var answer = model.chat(messages);  
  
        System.out.println(answer.aiMessage().text());  
    }  
}
```

Simple Completion with System Message

```
public class SimpleChatCompletion {  
    public static void main(String[] argv) {  
  
        ChatModel cmodel = AnthropicChatModel.builder()  
            .apiKey(System.getenv("ANTHROPIC_API_KEY"))  
            .modelName("CLAUDE_3_5 SONNET_20241022")  
            .timeout(Duration.ofSeconds(120))  
            .maxTokens(256)  
            .build();  
  
        List<ChatMessage> messages = new ArrayList<>();  
  
        SystemMessage sysmsg = new SystemMessage(""  
            You are a polite Java expert explaining concepts to a grammar school child.  
        "");  
        messages.add(sysmsg);  
  
        UserMessage usrmsg = UserMessage.from("Why should I learn Java.");  
        messages.add(usrmsg);  
  
        var answer = cmodel.chat(messages);  
  
        System.out.println(answer.aiMessage().text());  
    }  
}
```

ANTHROPIC



Check version!

```
def lc4jVersion = "1.4.0"
...
implementation "dev.langchain4j:langchain4j:${lc4jVersion}"
implementation "dev.langchain4j:langchain4j-open-ai:${lc4jVersion}"
implementation "dev.langchain4j:langchain4j-google-ai-gemini:${lc4jVersion}"
implementation "dev.langchain4j:langchain4j-anthropic:${lc4jVersion}"
implementation "dev.langchain4j:langchain4j-mistral-ai:${lc4jVersion}"
..."
```

Gradle dependencies
build.gradle

Use the LC4J libraries that you are using in your app

```
<dependency>
  <groupId>dev.langchain4j</groupId>
  <artifactId>langchain4j-core</artifactId>
  <version>1.4.0</version>
</dependency>

<dependency>
  <groupId>dev.langchain4j</groupId>
  <artifactId>langchain4j-open-ai</artifactId>
  <version>1.4.0</version>
</dependency>
```

Check version!

Maven dependencies
pom.xml

Use the LC4J libraries that you are using in your app

EMBEDDINGS DEMOS

©Frank Greco

```
public class GetEmbedding {
    public static void main(String[] argv) {
        String apiKey = System.getenv("OPENAI_API_KEY");

        EmbeddingModel model = OpenAiEmbeddingModel.builder(
            .apiKey(apiKey)
            .modelName(TEXT_EMBEDDING_3_SMALL)
            .build();

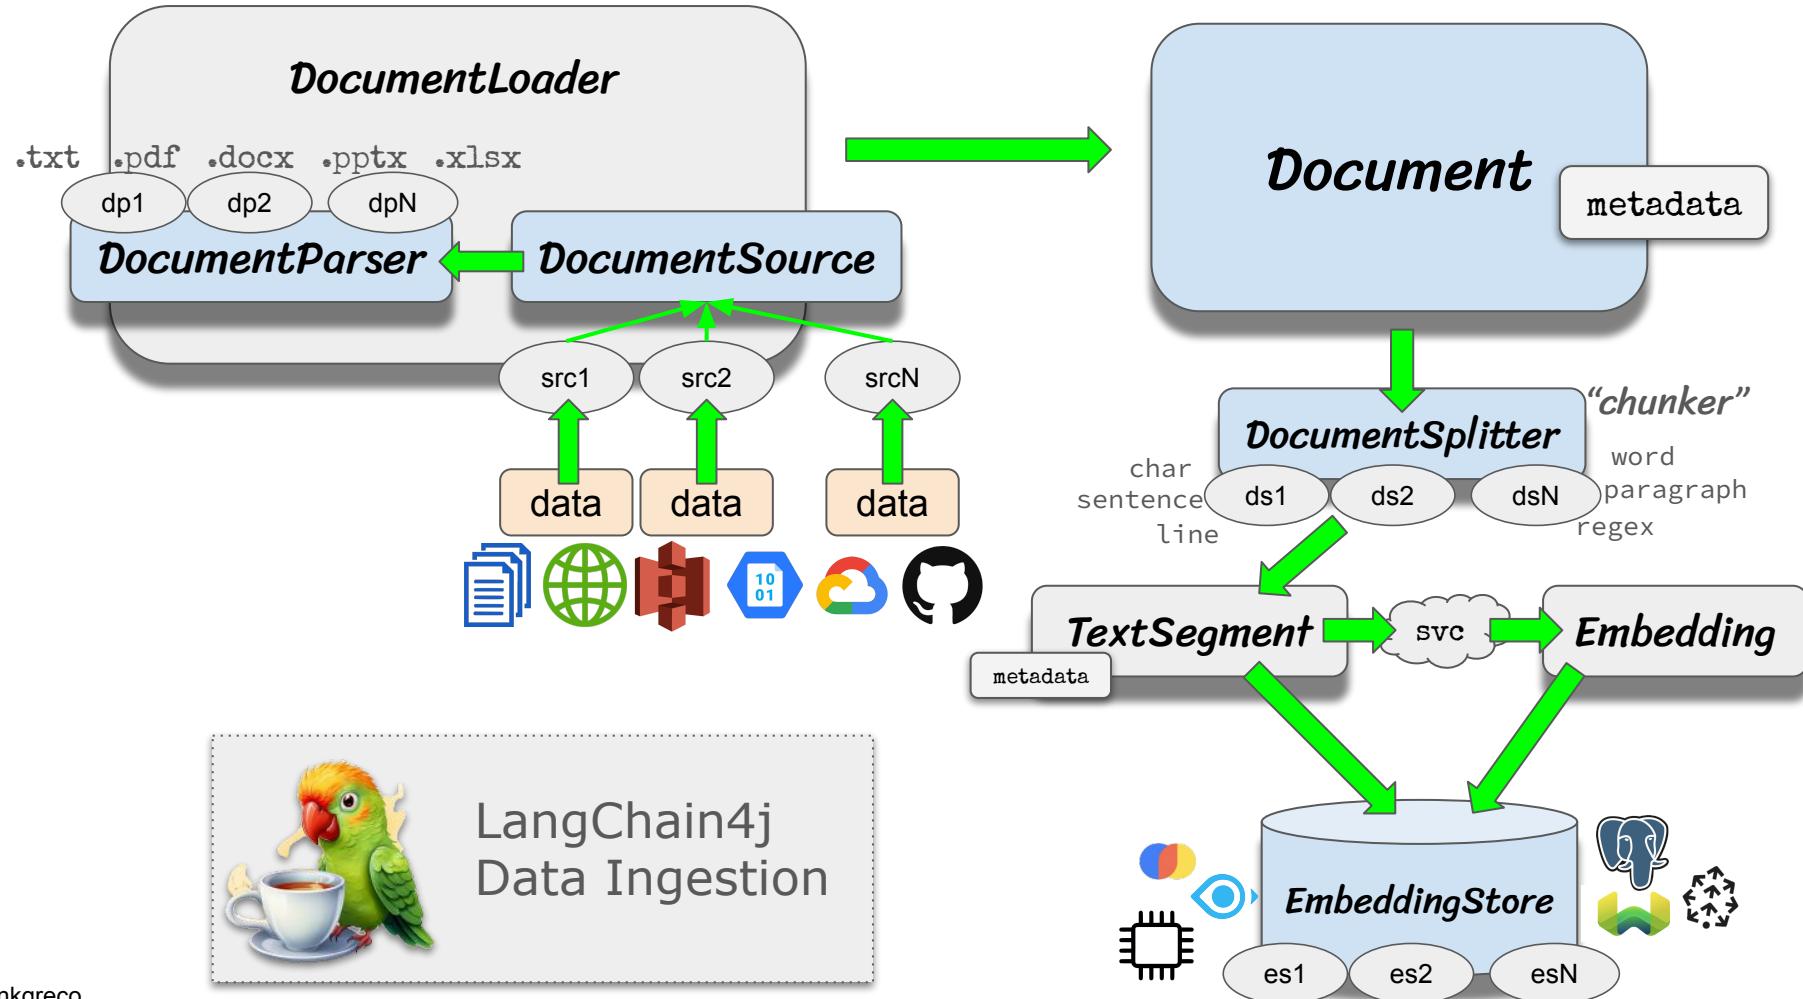
        Response<Embedding> response = model.embed(
            "One fish, two fish, red fish, blue fish");
        Embedding embedding = response.content();

        System.out.println(embedding.dimension() + " elements: " + embedding);
    }
}
```

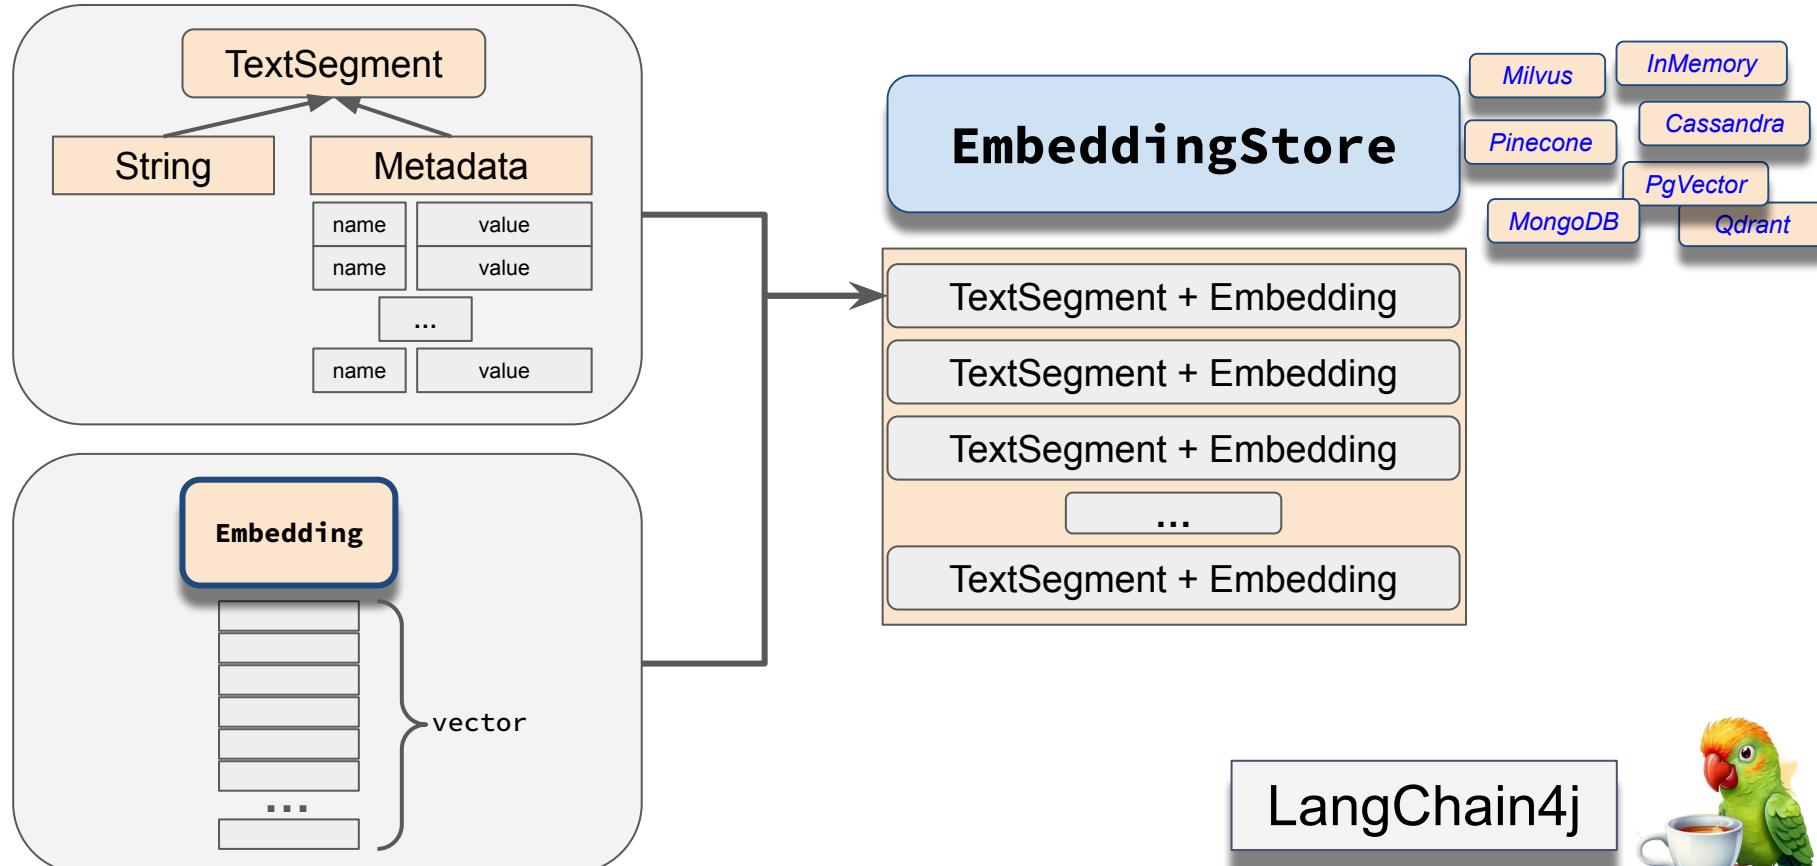
1536 elements: Embedding { vector = [-0.0052022343, -0.013655049,
0.018811593, 0.024425045, 0.024098681, -0.0039979527, -0.037362095,
-0.0030776078, -0.03321075, 3.4002998E-4, 0.06532492, -0.04566478,
0.002183372, 0.0138378125, 0.005965925, 0.010939704, -0.0035443075,
-0.0379626, -0.021461664, 0.051121578, 0.017245047, -0.06391503, 0.01575683,
0.022492973, 0.016827302, 0.032166388, 5.6379294E-4, 0.031043695,
0.021526936, 0.018876866, 0.07373204, -0.016970903, 0.0035475714,
-0.01642261, -0.057231102, 0.0049770433, 0.007206106, -0.02438588,
0.01518243, -0.0010810791, 0.018902974, 0.048641216, 0.05482907,
-0.013641994, 0.06960681, 0.020769773, -0.043314964, -0.028693879,
-0.005858225, 0.06558601, -0.040103547, 0.085585564, -0.054881286,
0.037623186, -0.0023171809, -0.008648633, -0.049528927, 0.03780595,
0.024059517, -0.058223248, 0.05477685, 0.032531913, 0.003802135,
0.0100128325, -0.020247592, 0.040991254, -0.047701288, -0.0178325,
0.052061506, 0.018459119, -0.016096247, 0.020456465, 0.014568866,
0.013093703, 0.04490762, ...

LANGCHAIN4J DATA INGESTION

©Frank Greco



Storing Embedding Vectors, Strings, and Metadata



VECTOR STORES AND CHUNKING

©Frank Greco

While you can implement this yourself
for educational purposes, use a library.

Most (embedding) vector databases offer a variety of
high-performance similarity functions.

Vector Database (VDB)

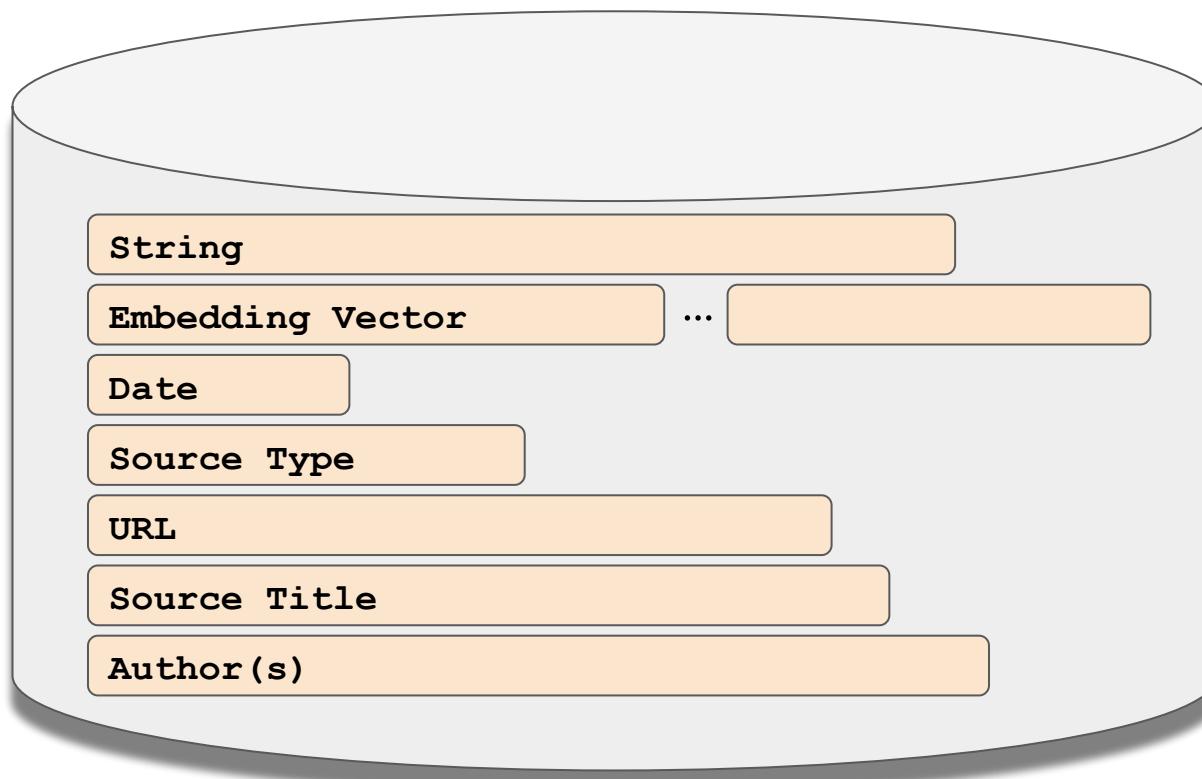
Stores the embedding vector and the original data.

Provides fast similarity search across all the vectors.

Can store metadata (the source URL of the original data, insert date, author, etc) for citations and other purposes.

Used for similarity searches, recommendation engines, anomaly detection, categorization, forecasting, and... RAG!

Vector Database (VDB)



Lots of VDBs



Milvus



redis



Chroma



Weaviate



DATASTAX

activeloop
Deep Lake



Vespa

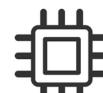


LanceDB

[ROCKSET]

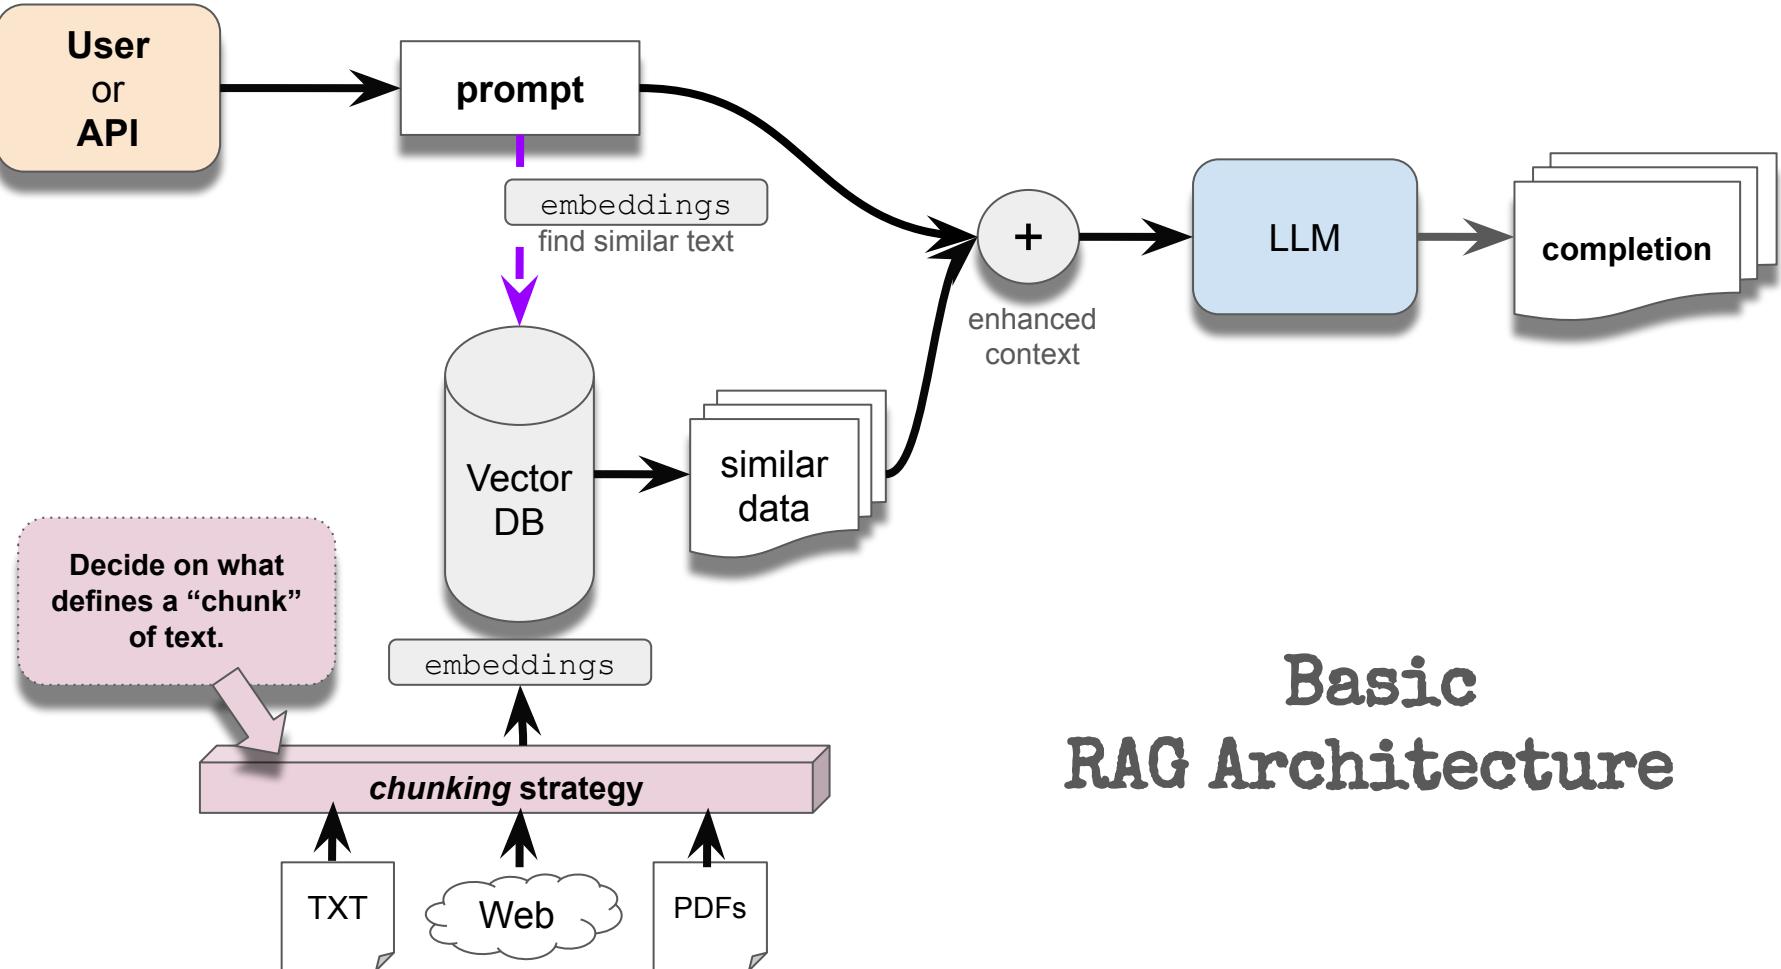


PostgreSQL



pgvector

drant



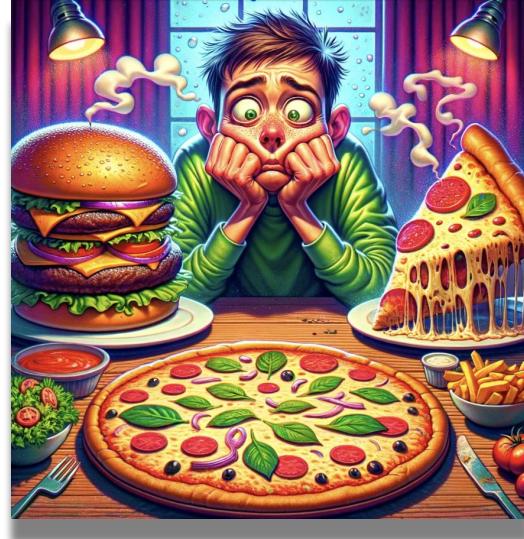
We need a “Chunking” Strategy

*A RAG system needs to know about **relevant** text to provide context in the prompt sent to the LLM.*

*This helps to improve **accuracy** .*

Due to context-window constraints and costs, it is not always feasible to send huge “chunks” of text.

Character limit?
Entire documents?
Sentences?
Paragraphs?
Phrases?
Markdown? HTML? JSON?
Recursive doc summarization chunkers?



Smaller chunks may be good for granular queries.
Larger chunks may be good for more in-depth queries.
Need to monitor/evaluate!

Performance of LLMs

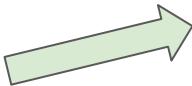
Chunking Strategies

Embedding Services

RAG Strategy

Vector Store

*These seem to have the
most positive effects*



LLM(s) Choices

LLM Parameters

SEARCHING AN EMBEDDINGSTORE

©Frank Greco

Searching the EmbeddingStore

LOW-LEVEL - more control, more complex

1. *Get embedding vector for target String*
2. *Create a EmbeddingSearchRequest with that vector*
3. *search(EmbeddingSearchRequest)*
4. *Retrieve matches*

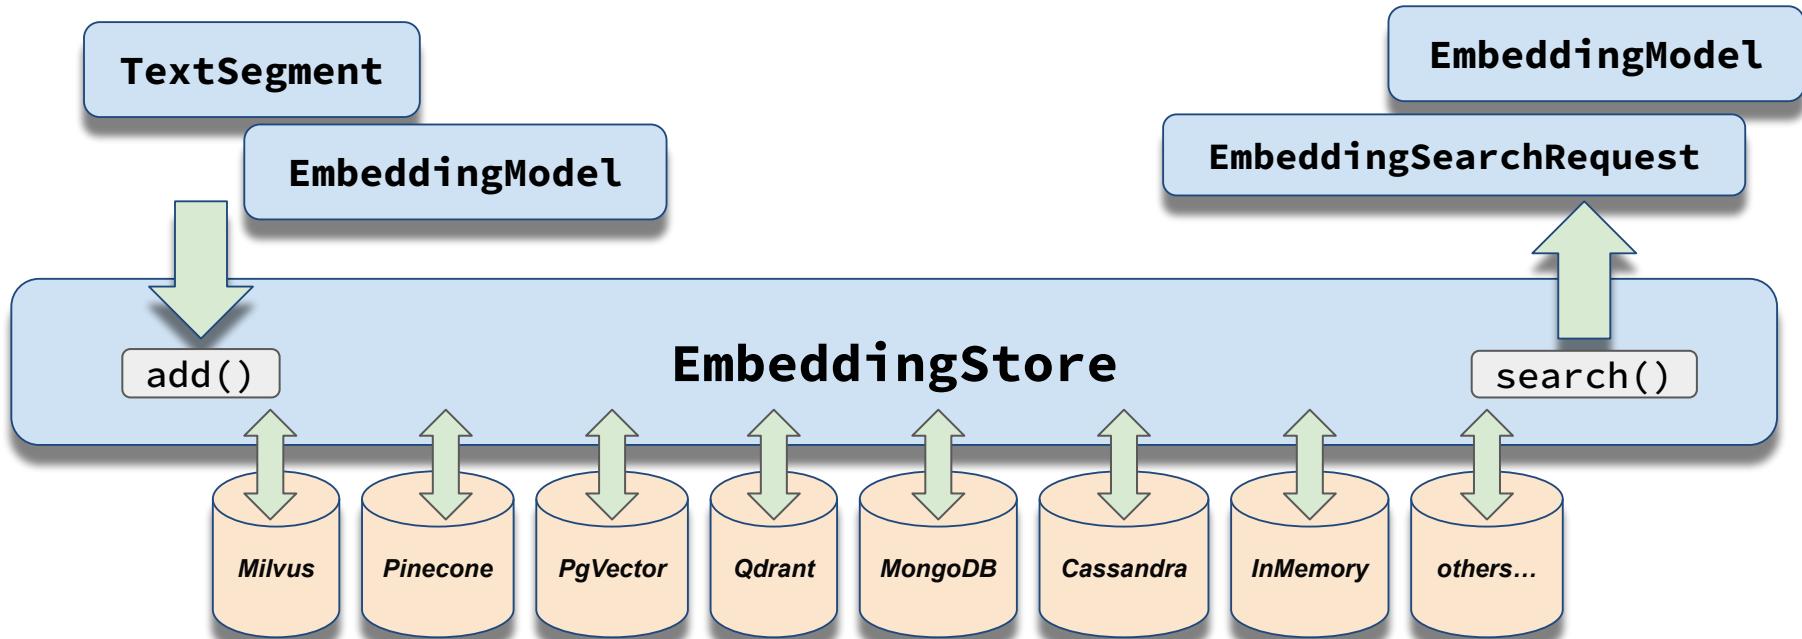
HIGH-LEVEL - easier

1. *Create EmbeddingStoreContentRetriever*
2. *Associate with AiService*

Using an EmbeddingStore

Low-Level

EmbeddingStore - Low-Level



LangChain4j



SEMANTIC SEARCH

DEMO

©Frank Greco

SemanticSearch – SemanticSearch.java [SemanticSearch.main]

SemanticSearch > src > main > java > search > SemanticSearch > main

SemanticSearch.java x Ingest.java

Project Commit Pull Requests

```
public class SemanticSearch {    Frank Greco
    public static void main(String[] args) throws IOException {        Frank Greco
        integer maxResults = integer.parseInt(prop.getProperty("SemanticSearch.maxResults"));
        Double minScore = Double.parseDouble(prop.getProperty("SemanticSearch.minScore"));
        boolean verbose = Boolean.parseBoolean(prop.getProperty("SemanticSearch.verbose"));

        while (true) {
            String query = console.readLine(pstring);
            if (set.contains(query.toLowerCase()))
                break;

            queryEmbedding = emodel.embed(query).content();
            EmbeddingSearchRequest embeddingSearchRequest = EmbeddingSearchRequest.builder()
                .queryEmbedding(queryEmbedding)
                .maxResults(maxResults)
                .minScore(minScore)
                .build();
            EmbeddingSearchResult<TextSegment> searchResult = myDB.search(embeddingSearchRequest);
            List<EmbeddingMatch<TextSegment>> matches = searchResult.matches();

            displayMatches(matches, verbose);
        }
    }
}

/**
 * displayMatches(matches, verbose)
 *
```

Run: SemanticSearch [:search.SemanticSearch.main()] x SemanticSearch [:database.Ingest.main()] x

SemanticSearch 23 sec 7:01:50 PM: Executing ':search.SemanticSearch.main()'
:search.Sema 22 sec

String>

Bookmarks Structure Bookmarks Structure

Git Run Debug Problems TODO Profiler Terminal Services Build

57:1 LF UTF-8 4 spaces main

SemanticSearch.java

Ingest.java



Thanks!

Questions / Comments?

LinkedIn Learning

Foundations of AI/ML for Java Developers

Introduction to Building GenAI Java Applications using LangChain4j

Pearson/O'Reilly

GenAI for Busy Java Developers

@frankgreco